

MSDS Practicum II - Diabetes Data Analysis

May 2, 2020

0.1 Sylvester Brown

0.2 Regis University

0.3 Praticum II

0.3.1 Spring 2020

0.3.2 Purpose of project:

Showcase skills in data manipulation and engineering, exploratory data analysis, visualizations and Machine leaning.

0.3.3 The problems to solve are:

1. What factors are best for predicting that someone has diabetes?
2. Which machine learning tool is best for predicting the outcome?

0.3.4 About the Dataset:

The dataset was contributed to Kaggle.com by Rahul. Data can be found at “<https://www.kaggle.com/rahulsah06/machine-learning-for-diabetes-with-python>”

0.3.5 Other Resources used:

”Pandas for Everyone’ by Daniel Chen

<https://www.kaggle.com/sagira/logistic-regression-math-behind-without-sklearn>

<https://www.kaggle.com/kwisatzhaderach/simple-mlp-approach>

<https://www.kaggle.com/aljojose/pima-indians-diabetes-my-first-project>

<https://towardsdatascience.com/feature-selection-techniques-in-machine-learning-with-python-f24e7da3f36e>

<https://www.kaggle.com/saurav9786/eda-makes-sense>

<https://www.kaggle.com/uciml/pima-indians-diabetes-database>

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger database. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

0.4 Variables

Pregnancies - Number of times pregnant

Glucose - Plasma glucose concentration a 2 hours in an oral glucose tolerance test

BloodPressure - Diastolic blood pressure (mm Hg)

SkinThickness - Triceps skin fold thickness (mm)

Insulin - 2-Hour serum insulin (mu U/ml)

BMI - Body mass index (weight in kg/(height in m)²)

DiabetesPedigreeFunction - Diabetes pedigree function

Age - Age (years)

Outcome

0.4.1 Import Libraries

```
[1]: # import numpy, pandas, seaborn and matplotlib

import numpy as np
import pandas as pd
import seaborn as sns
sns.set()
import matplotlib.pyplot as plt
%matplotlib inline

# import warnings filter
import warnings
warnings.filterwarnings('ignore')
```

0.4.2 Upload Data

```
[2]: # Use pandas to read the csv file

data= pd.read_csv('C:/Users/spbro/OneDrive/Desktop/diabetes_data.csv')
```

0.4.3 Exploratory Data Analysis

```
[3]: # List the column names for the dataset

data.columns
```

```
[3]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
          'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
          dtype='object')
```

```
[4]: # drop DiabetesPedigreeFunction column. Description does not appear to
# indicate what it really means and the purpose in the analysis

data.drop(['DiabetesPedigreeFunction'],axis=1,inplace=True)
```

```
[5]: # use rename() to change column names

data = data.rename(columns = {'Pregnancies':'pregnancies'})
data = data.rename(columns = {'Glucose':'glucose'})
data = data.rename(columns = {'BloodPressure':'bloodpressure'})
data = data.rename(columns = {'Insulin':'insulin'})
data = data.rename(columns = {'Age':'age'})
data = data.rename(columns = {'Outcome':'outcome'})
data = data.rename(columns = {'SkinThickness':'skinthickness'})
data = data.rename(columns = {'BMI':'bmi'})
```

```
[6]: # Display first 5 rows

data.head()
```

```
[6]:   pregnancies  glucose  bloodpressure  skinthickness  insulin   bmi  age  \
0             6     148             72             35         0  33.6  50
1             1      85             66             29         0  26.6  31
2             8     183             64              0         0  23.3  32
3             1      89             66             23        94  28.1  21
4             0     137             40             35       168  43.1  33

      outcome
0           1
1           0
2           1
3           0
4           1
```

```
[7]: # display columns

data.columns
```

```
[7]: Index(['pregnancies', 'glucose', 'bloodpressure', 'skinthickness', 'insulin',
          'bmi', 'age', 'outcome'],
          dtype='object')
```

```
[8]: # Display last 5 rows
```

```
data.tail()
```

```
[8]:      pregnancies  glucose  bloodpressure  skinthickness  insulin  bmi  age  \
763             10     101             76             48     180  32.9  63
764              2     122             70             27       0  36.8  27
765              5     121             72             23    112  26.2  30
766              1     126             60              0       0  30.1  47
767              1      93             70             31       0  30.4  23

      outcome
763         0
764         0
765         0
766         1
767         0
```

```
[9]: data.shape
```

```
# There are 768 rows and 9 columns
```

```
[9]: (768, 8)
```

```
[10]: # A statistical description of each column
```

```
data.describe()
```

```
[10]:      pregnancies      glucose  bloodpressure  skinthickness      insulin  \
count    768.000000    768.000000    768.000000    768.000000    768.000000
mean       3.845052    120.894531     69.105469     20.536458     79.799479
std        3.369578     31.972618     19.355807     15.952218    115.244002
min         0.000000     0.000000     0.000000     0.000000     0.000000
25%         1.000000     99.000000     62.000000     0.000000     0.000000
50%         3.000000    117.000000     72.000000     23.000000     30.500000
75%         6.000000    140.250000     80.000000     32.000000    127.250000
max        17.000000    199.000000    122.000000     99.000000    846.000000

      bmi      age      outcome
count    768.000000    768.000000    768.000000
mean     31.992578     33.240885     0.348958
std       7.884160     11.760232     0.476951
min        0.000000     21.000000     0.000000
25%       27.300000     24.000000     0.000000
50%       32.000000     29.000000     0.000000
75%       36.600000     41.000000     1.000000
max       67.100000     81.000000     1.000000
```

```
[11]: # check if data contains any null values

data.isnull().any()
```

```
[11]: pregnancies      False
      glucose          False
      bloodpressure    False
      skinthickness    False
      insulin          False
      bmi              False
      age              False
      outcome          False
      dtype: bool
```

```
[12]: # display more descriptive stats using .info()

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   pregnancies           768 non-null   int64  
 1   glucose               768 non-null   int64  
 2   bloodpressure         768 non-null   int64  
 3   skinthickness         768 non-null   int64  
 4   insulin               768 non-null   int64  
 5   bmi                   768 non-null   float64 
 6   age                   768 non-null   int64  
 7   outcome               768 non-null   int64  
dtypes: float64(1), int64(7)
memory usage: 48.1 KB
```

```
[13]: # how many instances of diabetes using count

data['outcome'].value_counts()
```

```
[13]: 0      500
      1      268
      Name: outcome, dtype: int64
```

```
[14]: # Use groupby to explore number of people with diabetes per each observation,
      ↪ using mean values

data.groupby('outcome').mean()
```

```
[14]:      pregnancies      glucose  bloodpressure  skinthickness      insulin  \
outcome
0          3.298000   109.980000      68.184000      19.664000    68.792000
1          4.865672   141.257463      70.824627      22.164179   100.335821

      bmi      age
outcome
0      30.304200  31.190000
1      35.142537  37.067164
```

Although the data info shows that there are no missing values, I can tell that the data is skewed because of unusual number of zeros. It appears unusual that Glucose, BloodPressure, SkinThickness, Insulin, and BMI will have zeros. I will concentrate on these five attributes and replace the zeros with NaN and then replace with the mean values.

```
[15]: #Replace 0 with NaN

d=data[['glucose','bloodpressure','skinthickness','insulin','bmi']]=data[['glucose','bloodpres
↪replace(0,np.nan)
d.head()
```

```
[15]:      glucose  bloodpressure  skinthickness  insulin  bmi
0      148.0          72.0          35.0      NaN  33.6
1      85.0          66.0          29.0      NaN  26.6
2      183.0          64.0          NaN      NaN  23.3
3      89.0          66.0          23.0     94.0  28.1
4      137.0          40.0          35.0    168.0  43.1
```

```
[16]: #Replace NaN with mean values

data['glucose'].fillna(data['glucose'].median(),inplace=True)
data['bloodpressure'].fillna(data['bloodpressure'].median(),inplace=True)
data['skinthickness'].fillna(data['skinthickness'].median(),inplace=True)
data['insulin'].fillna(data['insulin'].median(),inplace=True)
data['bmi'].fillna(data['bmi'].median(),inplace=True)
data.head()
```

```
[16]:      pregnancies  glucose  bloodpressure  skinthickness  insulin  bmi  age  \
0          6      148.0          72.0          35.0    125.0  33.6   50
1          1      85.0          66.0          29.0    125.0  26.6   31
2          8      183.0          64.0          29.0    125.0  23.3   32
3          1      89.0          66.0          23.0     94.0  28.1   21
4          0      137.0          40.0          35.0    168.0  43.1   33

      outcome
0          1
1          0
```

2	1
3	0
4	1

0.5 Data Visualization

0.6 BAR PLOT

```
[17]: # The plot shows the amount of individuals that have diabetes and those who
      ↪ dont.

      # countplot---Plot the frequency of the Outcome

      fig1, ax1 = plt.subplots(1,2,figsize=(8,8))

      #It shows the count of observations in each categorical bin using bars


      sns.countplot(data['outcome'],ax=ax1[0])

      #Find the % of diabetic and Healthy person

      labels = 'Healthy', 'Diabetic'

      data.outcome.value_counts().plot.pie(labels=labels, autopct='%1.
      ↪1f%%',shadow=True, startangle=90)
```

```
[17]: <matplotlib.axes._subplots.AxesSubplot at 0x23bf994b188>
```



0.6.1 BOXPLOTS

```
[18]: f, axes = plt.subplots(4,2, figsize=(10,10), sharex=True)

plt.subplots_adjust(wspace=0.5) # adjust the space between the plots

sns.despine(left=True)

# plot a boxplot of each observation to see if there are outliers

sns.boxplot( x= 'pregnancies', data=data, orient='v',ax=axes[0,0])

sns.boxplot( x= 'glucose', data=data, orient='v',ax=axes[0,1])

sns.boxplot( x= 'bloodpressure', data=data, orient='v',ax=axes[1,0])
```



```

sns.boxplot( x= 'skinthickness', data=data, orient='v',ax=axes[1,1])

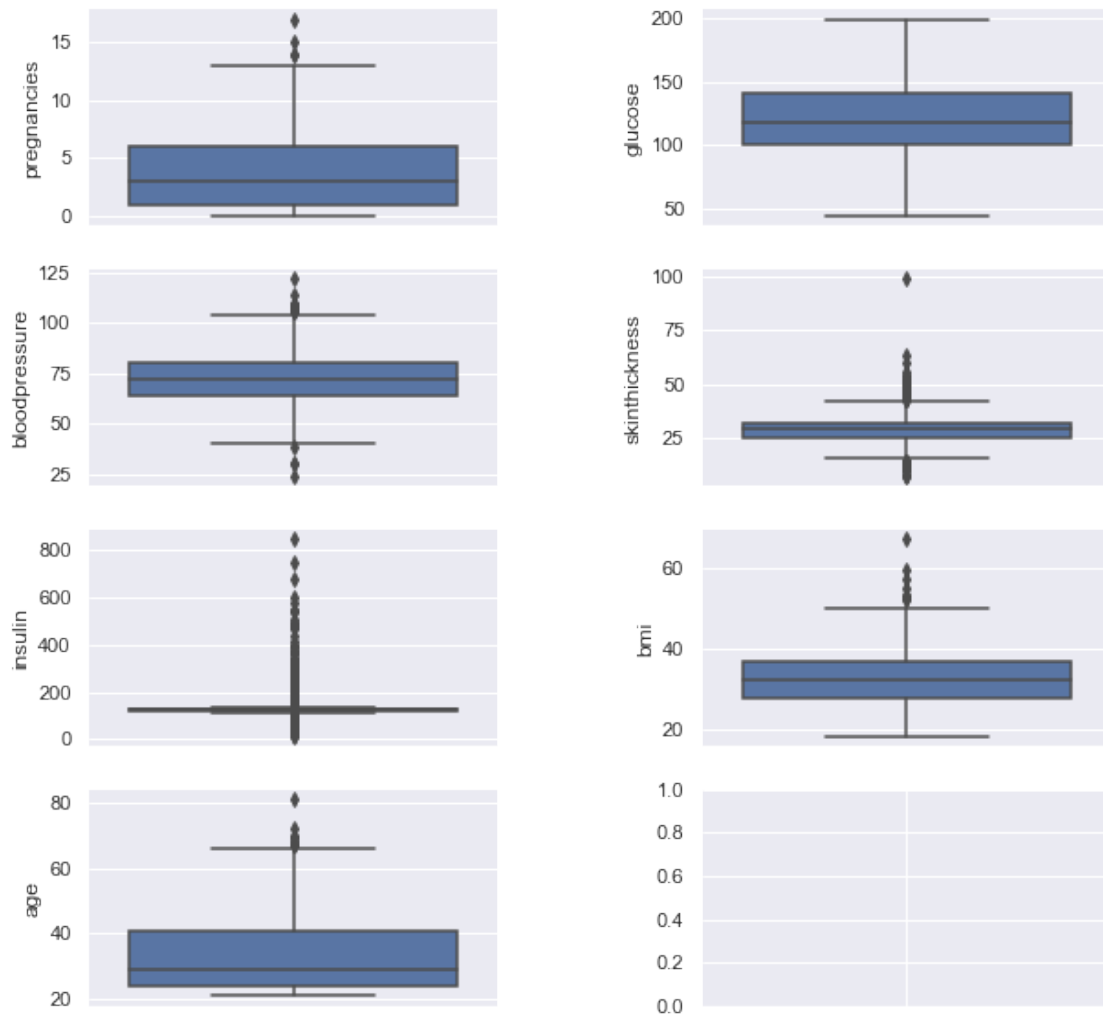
sns.boxplot( x= 'insulin', data=data, orient='v',ax=axes[2,0])

sns.boxplot( x= 'bmi', data=data, orient='v',ax=axes[2,1])

sns.boxplot( x= 'age', data=data, orient='v',ax=axes[3,0])

```

[18]: <matplotlib.axes._subplots.AxesSubplot at 0x23bf9b5f048>

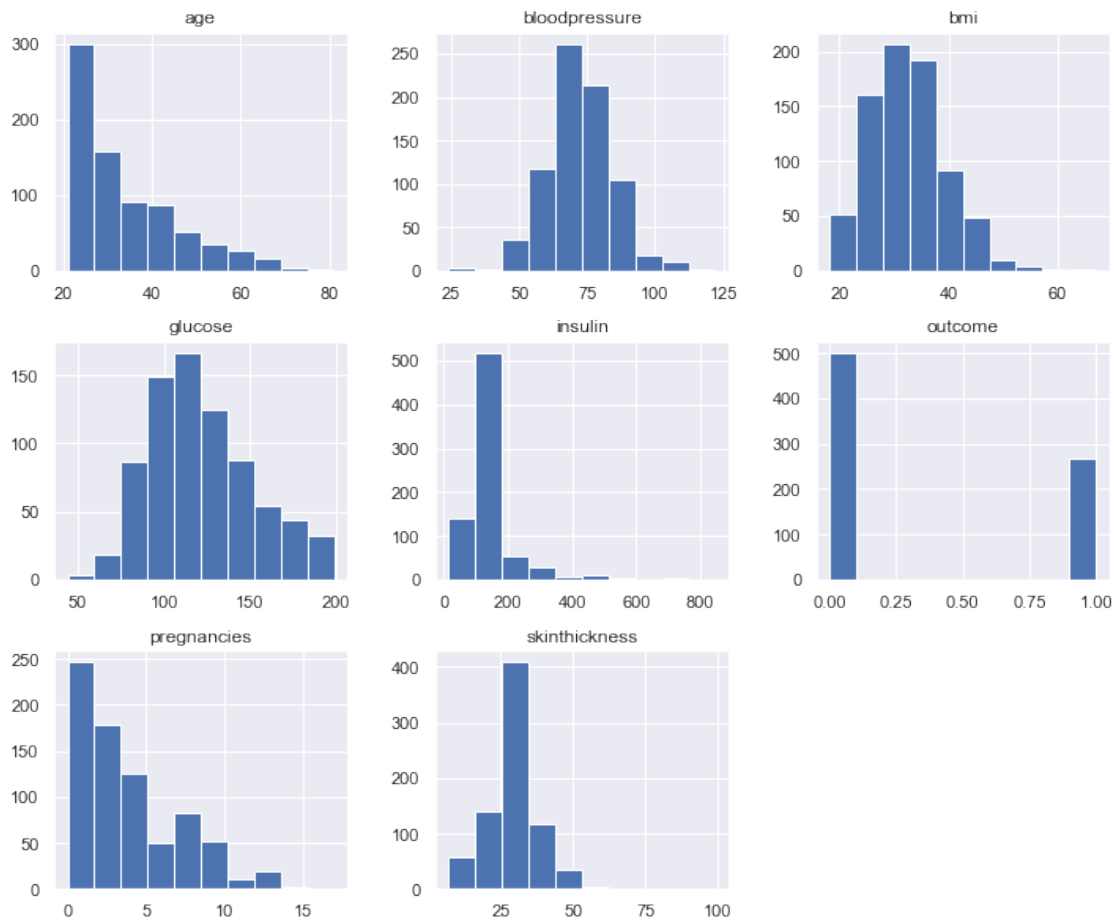


0.7 Observation

Although we have replaced zeros, the boxplots show outliers in all attributes.

0.8 HISTOGRAM

```
[19]: data.hist(figsize=(12,10));
```



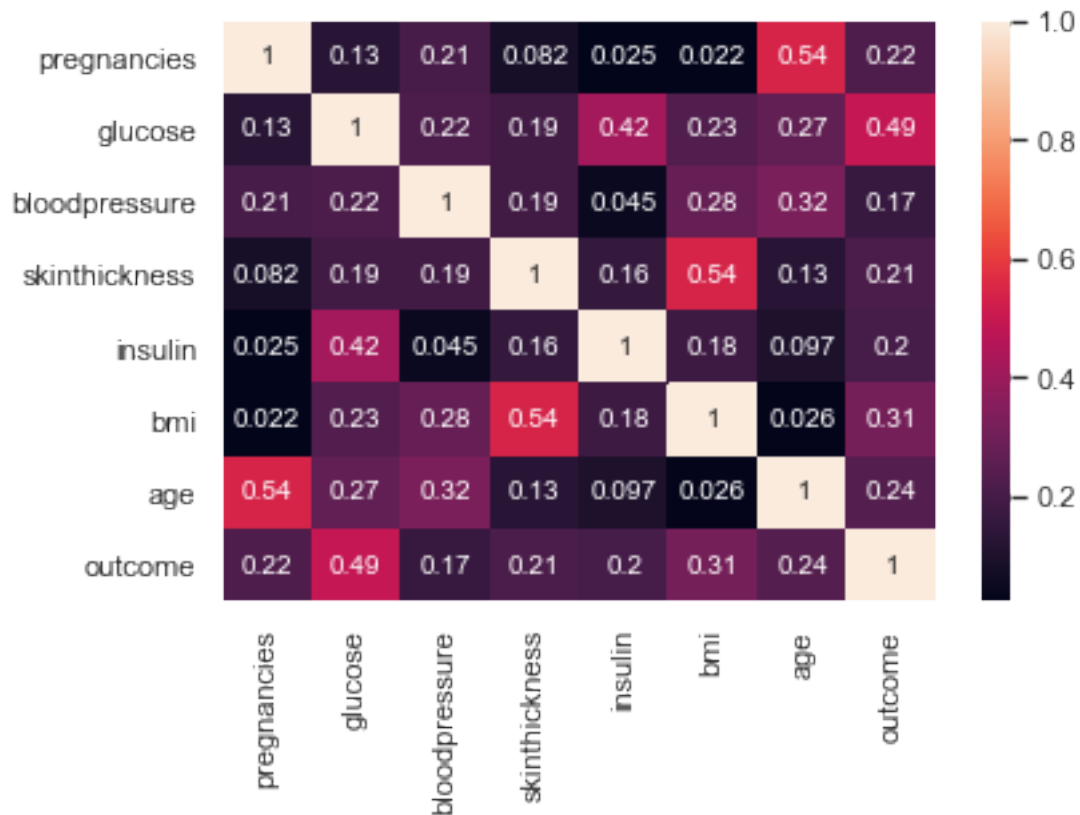
0.9 OBSERVATION

The plots show that Glucose, Blood Pressure, BMI are normally distributed. Pregnancies, Insulin, and Age are right-skewed.

0.10 CORRELATION PLOT

```
[20]: # correlation matrix  
  
cor=data.corr()  
  
# correlation plot---heatmap  
  
sns.heatmap(cor,annot=True)
```

[20]: <matplotlib.axes._subplots.AxesSubplot at 0x23bfa521e08>



0.11 OBSERVATION

This shows that glucose is the most correlated with the outcome. Also, bmi, insulin and age appear to be next on the correlation scale.

Other correlation pairs are pregnancy to age, insulin to glucose, bmi and skinthickness, age and glucose.

0.12 MACHINE LEARNING (ML)

[21]: *# Split data into X = observation and y = outcome*

```
attributes = list(data.columns[:7])
X = data[attributes].values
y= data['outcome'].values

from sklearn import preprocessing

# Get column names first
```


```
names = data.columns[:7]# Create the Scaler object
scaler = preprocessing.StandardScaler()# Fit your data on the scaler object
scaled_df = scaler.fit_transform(X)
X = pd.DataFrame(scaled_df, columns=names)
```

Use ExtraTreesClassifier to determine feature or observation importance

```
[22]: from sklearn.ensemble import ExtraTreesClassifier
model = ExtraTreesClassifier()
model.fit(X,y)
print(model.feature_importances_)

#plot graph of feature importances for better visualization
feat_importances = pd.Series(model.feature_importances_, index=X.columns)
feat_importances.nlargest(10).plot(kind='barh')
plt.show()
```

```
[0.10963028 0.28086509 0.10084644 0.10237593 0.09284294 0.16765796
0.14578137]
```



0.13 OBSERVATION

It was surprising to me based on research that insulin and blood pressure will be lower on the scale to predict the outcome. Pregnancy was a bit more of a preferred predictor of the outcome than blood pressure. The attribute that has a better predictability of the outcome was glucose.

Given that most of the observations were skewed, I decided to standardize the dataset in preparation for machine learning. This will be simply achieved using StandardScaler from Sklearn.preprocessing.

```
[23]: # Split into train and test sets for machine learning.
# test set was split at 20% and random_state set to 42

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20,
    random_state = 42)
```

0.14 MODELS for ML

```
[24]: # Import suite of algorithms.
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis,
    QuadraticDiscriminantAnalysis
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier,
    BaggingClassifier, GradientBoostingClassifier
from sklearn.calibration import CalibratedClassifierCV
```

```
[25]: # Create objects of required models.
models = []
models.append(("KNN", KNeighborsClassifier()))
models.append(("DecisionTree", DecisionTreeClassifier()))
models.append(("QDA", QuadraticDiscriminantAnalysis()))
models.append(("AdaBoost", AdaBoostClassifier()))
models.append(("SVM Linear", SVC(kernel="linear")))
models.append(("Random Forest", RandomForestClassifier()))
models.append(("Bagging", BaggingClassifier()))
```

```
[26]: # Find accuracy of models.
results = []
for name, model in models:
    kfold = KFold(n_splits=10, random_state=123)
    cv_result = cross_val_score(model, X_train, y_train, cv = kfold, scoring =
    accuracy)
    results.append(tuple([name, cv_result.mean(), cv_result.std()]))

results.sort(key=lambda x: x[1], reverse = True)
for i in range(len(results)):
    print('{:20s} {:.2f} (+/-) {:.2f} '.format(results[i][0], results[i][1],
    * 100, results[i][2] * 100))
```

SVM Linear	75.41 (+/-) 4.19
Bagging	75.09 (+/-) 2.93
AdaBoost	75.08 (+/-) 4.48
Random Forest	73.95 (+/-) 3.48
QDA	73.78 (+/-) 6.01
KNN	73.30 (+/-) 6.82
DecisionTree	71.99 (+/-) 4.22

0.15 FIND BEST PARAMETERS TO OPTIMIZE BEST MODEL

```
[27]: ### SVM Linear was the best model.
      ### find the parameters for SVM that will optimize the model
      ### using GridSearchCV.

from sklearn.model_selection import GridSearchCV
model = SVC()
paramaters = [
    {'C' : [0.001, 0.01, 0.1, 1, 10, 100, 1000], 'kernel' : ['linear']}
]
grid_search = GridSearchCV(estimator = model,
                           param_grid = paramaters,
                           scoring = 'accuracy',
                           cv = 10,
                           n_jobs = -1)
grid_search = grid_search.fit(X_train, y_train)
best_accuracy = grid_search.best_score_
best_parameters = grid_search.best_params_

print('Best accuracy : ', grid_search.best_score_)
print('Best parameters :', grid_search.best_params_ )
```

Best accuracy : 0.758957654723127

Best parameters : {'C': 0.01, 'kernel': 'linear'}

C = It controls the trade off between smooth decision boundary and classifying training points correctly. A large value of c means you will get more training points correctly.

Kernel = type of hyperplane used to separate the data

0.16 OPTIMIZE MODEL

0.16.1 Use the Best parameters : {'C': 0.01, 'kernel': 'linear'} to optimize the model

```
[28]: # Predict output for test set.
final_model = SVC(C = 0.01, kernel = 'linear')
final_model.fit(X_train, y_train)
y_pred = final_model.predict(X_test)
```

```

from sklearn.metrics import confusion_matrix, accuracy_score
cf = confusion_matrix(y_test, y_pred)
print(cf)
print(accuracy_score(y_test, y_pred) * 100)

from sklearn.metrics import classification_report
report = classification_report(y_test, y_pred)
print(report)

```

```

[[87 12]
 [23 32]]
77.27272727272727

```

	precision	recall	f1-score	support
0	0.79	0.88	0.83	99
1	0.73	0.58	0.65	55
accuracy			0.77	154
macro avg	0.76	0.73	0.74	154
weighted avg	0.77	0.77	0.77	154

0.16.2 Machine learning with Added Column - random numbers

```

[29]: # copy and modify original dataset for analysis
data2 = data

np.random.seed(42)

# create new column with the random number generator

data2['randNumCol'] = np.random.randint(10,40, size=len(data2))

```

```

[30]: # Rearrange columns so it is easier to split X and Y for analysis

data2=data2[['pregnancies', 'glucose', 'bloodpressure', 'skinthickness',
→'insulin',
      'bmi', 'age', 'randNumCol', 'outcome']]

```

```

[31]: # display first five rows
data2.head()

```

```

[31]:  pregnancies  glucose  bloodpressure  skinthickness  insulin  bmi  age  \
0           6    148.0         72.0         35.0    125.0  33.6  50
1           1     85.0         66.0         29.0    125.0  26.6  31
2           8    183.0         64.0         29.0    125.0  23.3  32
3           1     89.0         66.0         23.0     94.0  28.1  21

```

4	0	137.0	40.0	35.0	168.0	43.1	33
---	---	-------	------	------	-------	------	----

	randNumCol	outcome
0	16	1
1	29	0
2	38	1
3	24	0
4	20	1

```
[32]: # Split new data into X = observation and y = outcome

attributes = list(data2.columns[:8])
X = data2[attributes].values
y= data2['outcome'].values

from sklearn import preprocessing


# Get column names first
names = data2.columns[:8]# Create the Scaler object
scaler = preprocessing.StandardScaler()# Fit your data on the scaler object
scaled_df = scaler.fit_transform(X)
X = pd.DataFrame(scaled_df, columns=names)
```

```
[33]: # Use ExtraTreesClassifier to determine order of importance for attributes.

from sklearn.ensemble import ExtraTreesClassifier
model = ExtraTreesClassifier()
model.fit(X,y)
print(model.feature_importances_)

#plot graph of feature importances for better visualization
feat_importances = pd.Series(model.feature_importances_, index=X.columns)
feat_importances.nlargest(10).plot(kind='barh')
plt.show()
```

```
[0.10138801 0.25342595 0.09905497 0.08679671 0.06879194 0.1451958
 0.15091407 0.09443254]
```

0.16.3 OBSERVATION

The attributes that show up after the random number attribute shows the cut of the feature importance. This appears to indicate that the random number attribute is a better predictor of the outcome than the attributes that come after that.

```
[34]: # drop all attributes that placed below and including the random number column
```

```
X.drop(['insulin', 'skinthickness', 'randNumCol'],axis=1,inplace=True)
```

```
[35]: X.head()
```

```
[35]:
```

	pregnancies	glucose	bloodpressure	bmi	age
0	0.639947	0.866045	-0.031990	0.166619	1.425995
1	-0.844885	-1.205066	-0.528319	-0.852200	-0.190672
2	1.233880	2.016662	-0.693761	-1.332500	-0.105584
3	-0.844885	-1.073567	-0.528319	-0.633881	-1.041549
4	-1.141852	0.504422	-2.679076	1.549303	-0.020496

```
[36]: # Split feature importance attributes into train and test sets for machine
      ↪ learning.
      # test set was split at 20% and random_state set to 42

      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20,
      ↪ random_state = 42)
```

```
[37]: # Use Optimized model of ML to test prediction of the feature importance
      ↪ attributes.
      # Predict output for test set.
      final_model = SVC(C = 0.1, kernel = 'linear')
      final_model.fit(X_train, y_train)
      y_pred = final_model.predict(X_test)

      from sklearn.metrics import confusion_matrix, accuracy_score
      cf = confusion_matrix(y_test, y_pred)
      print(cf)
      print(accuracy_score(y_test, y_pred) * 100)

      from sklearn.metrics import classification_report
      report = classification_report(y_test, y_pred)
      print(report)
```

```
[[85 14]
 [20 35]]
77.92207792207793
```

	precision	recall	f1-score	support
0	0.81	0.86	0.83	99
1	0.71	0.64	0.67	55
accuracy			0.78	154
macro avg	0.76	0.75	0.75	154
weighted avg	0.78	0.78	0.78	154

0.16.4 Compare the two Predictions

The model improved after removing attributes that appear to not contribute much to the prediction. It increased from about 76% to 78%. Upon closer observation it appears that the improvement in the model was only reflected in prediction True negatives. Therefore it was not a significant change when the features or attributes were taken out. Further, given that the “C” parameter in SVM linear is used to determine how accurate a prediction should be versus the smoothness level of the classifications, a higher value of “C” provides stronger accuracy in predictions. This appears to explain why myb accuracy for the model is at about 78%.